Chicago Bears 2025 Season Prediction Data Science Capstone

Machine Learning • Logistic Regression • Predictive Analytics

By: Vyaas Subramanian

Project Overview

Goal of the Project

- Use machine learning to analyze and predict the Chicago Bears' performance for the 2025 season.
- Demonstrate a complete end-to-end data science workflow on real NFL data.

Why This Project Matters

- Sports analytics is one of the fastest growing areas in data science.
- Predictive modeling provides insights into team strengths, matchups, and season expectations.
- This project illustrates how data-driven methods can complement traditional sports analysis.

What This Analysis Includes

- Historical data extraction (2017–2024)
- Feature engineering
- Logistic regression training & evaluation
- Full-season win probability predictions
- Professional visualizations to communicate results

Data Cleaning & Preparation

DATASET:
NFL GAME RESULTS FROM 2017–2025 (REGULAR SEASON ONLY).

ACTIONS TAKEN:

- FILTERED THE DATASET TO ALL CHICAGO BEARS GAMES.
- REMOVED:
  - PRESEASON GAMES
  - ROWS WITH INVALID OPPONENTS OR MISSING SCORES
- STANDARDIZED TEAM NAMES TO ENSURE CONSISTENCY.
- CREATED THE BINARY OUTCOME VARIABLE:
- BEARS_WIN = 1 (WIN), 0 (LOSS)

WHY THIS MATTERS:
A CLEAN DATASET ENSURES THE MODEL LEARNS MEANINGFUL PATTERNS INSTEAD OF NOISE.

Features Used in the Model

FEATURE 1: ISHOME

- 1 = BEARS PLAYED AT HOME
- 0 = BEARS PLAYED AWAY

WHY IMPORTANT?

HOME-FIELD ADVANTAGE SIGNIFICANTLY AFFECTS NFL OUTCOMES DUE TO CROWD IMPACT, TRAVEL FATIGUE, AND ENVIRONMENTAL FAMILIARITY.

FEATURE 2: OPPONENT (ONE-HOT ENCODED)

- EACH OPPOSING TEAM BECOMES A BINARY INDICATOR (E.G., OPPONENT_PACKERS = 1 IF PLAYING PACKERS).

WHY IMPORTANT?

- OPPONENT STRENGTH IS ONE OF THE STRONGEST PREDICTORS OF GAME OUTCOMES.
- THIS ALLOWS THE MODEL TO LEARN WHICH MATCHUPS HISTORICALLY FAVOR OR HURT THE BEARS.

WHY ONLY THESE FEATURES?

- PRESENT FOR ALL SEASONS (NO MISSING DATA).
- AVOIDS OVERFITTING ON A SMALL DATASET.
- PRODUCES AN INTERPRETABLE MODEL SUITABLE FOR A MINI-CAPSTONE PROJECT.

Model Development

Chosen Model: Logistic Regression

Why this model?

- Ideal for binary prediction (win/loss)
- Highly interpretable coefficients
- Performs well on smaller datasets
- Easy to evaluate and visualize

Training Process:

- Used data from 2017–2024 only (2025 excluded for prediction).
- Split the dataset into:
  - 80% training
  - 20% testing

Model Output:

Probability the Bears win each game.

Logistic Regression

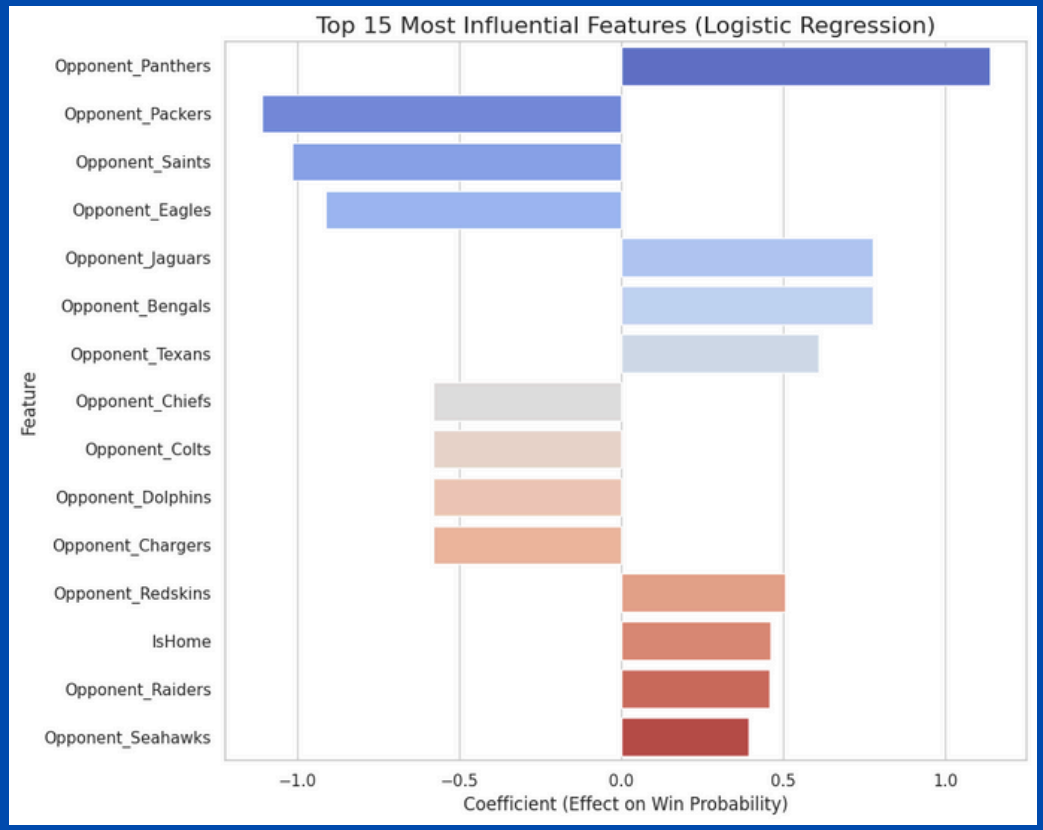

- SHOWS HOW EACH OPPONENT AND GAME LOCATION INFLUENCES WIN PROBABILITY.
- NEGATIVE COEFFICIENTS (E.G., PACKERS, SAINTS, EAGLES) INDICATE HISTORICALLY CHALLENGING OPPONENTS.
- POSITIVE COEFFICIENTS REFLECT OPPONENTS THE BEARS HAVE MATCHED UP BETTER AGAINST.
- THE ISHOME FEATURE HAS A MODERATE POSITIVE COEFFICIENT, CONFIRMING THE ADVANTAGE OF PLAYING AT SOLDIER FIELD.
- PROVIDES TRANSPARENCY INTO HOW THE MODEL MAKES PREDICTIONS CRUCIAL IN AN INTERPRETABLE MODEL LIKE LOGISTIC REGRESSION.

Model Evaluation Results

Performance Metrics (on test set):

- Accuracy: 71%
- Precision (Losses): 70%
- Recall (Losses): 94%
- Precision (Wins): 80%
- Recall (Wins): 36%

Interpretation:

- The model is very strong at identifying losses, reflecting Bears' performance patterns.
- The model is more conservative predicting wins, which is typical with limited features.
- Overall accuracy is solid given the small dataset and high unpredictability of sports.

Confusion Matrix

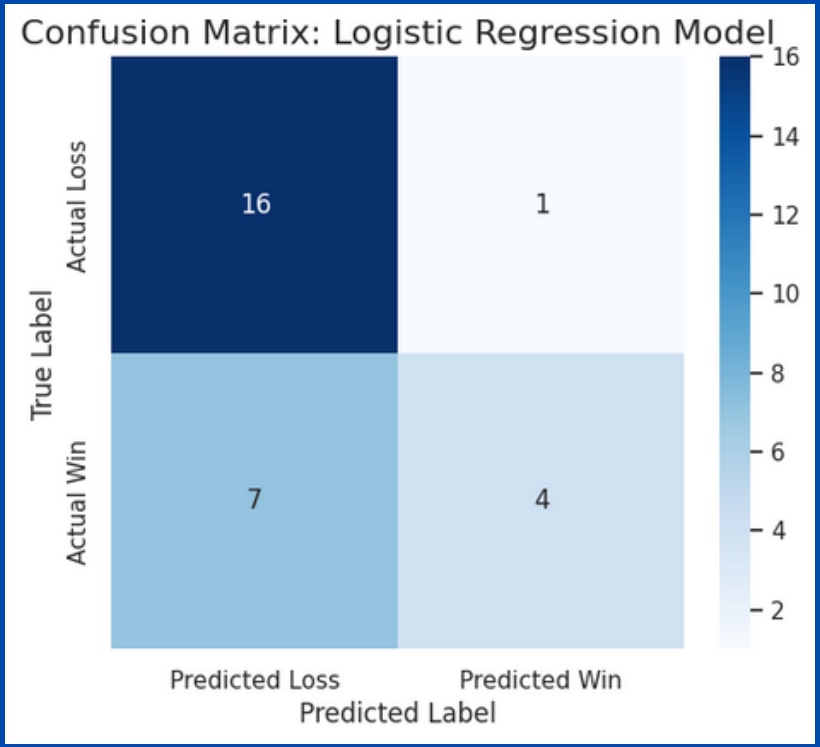

- HIGHLIGHTS HOW WELL THE MODEL PERFORMS ON HELD-OUT TEST DATA.
- THE MODEL ACCURATELY PREDICTS 16 OUT OF 17 LOSSES, SHOWING STRONG RELIABILITY IN IDENTIFYING UNFAVORABLE MATCHUPS.
- THE MODEL STRUGGLES WITH WINS (ONLY 4 CORRECT), WHICH IS COMMON WITH LIMITED FEATURES AND SMALL WIN COUNTS.
- OVERALL, THE MATRIX SHOWS THAT THE MODEL IS CONSERVATIVE AND LEANS TOWARD PREDICTING LOSSES.

ROC Curve (AUC)

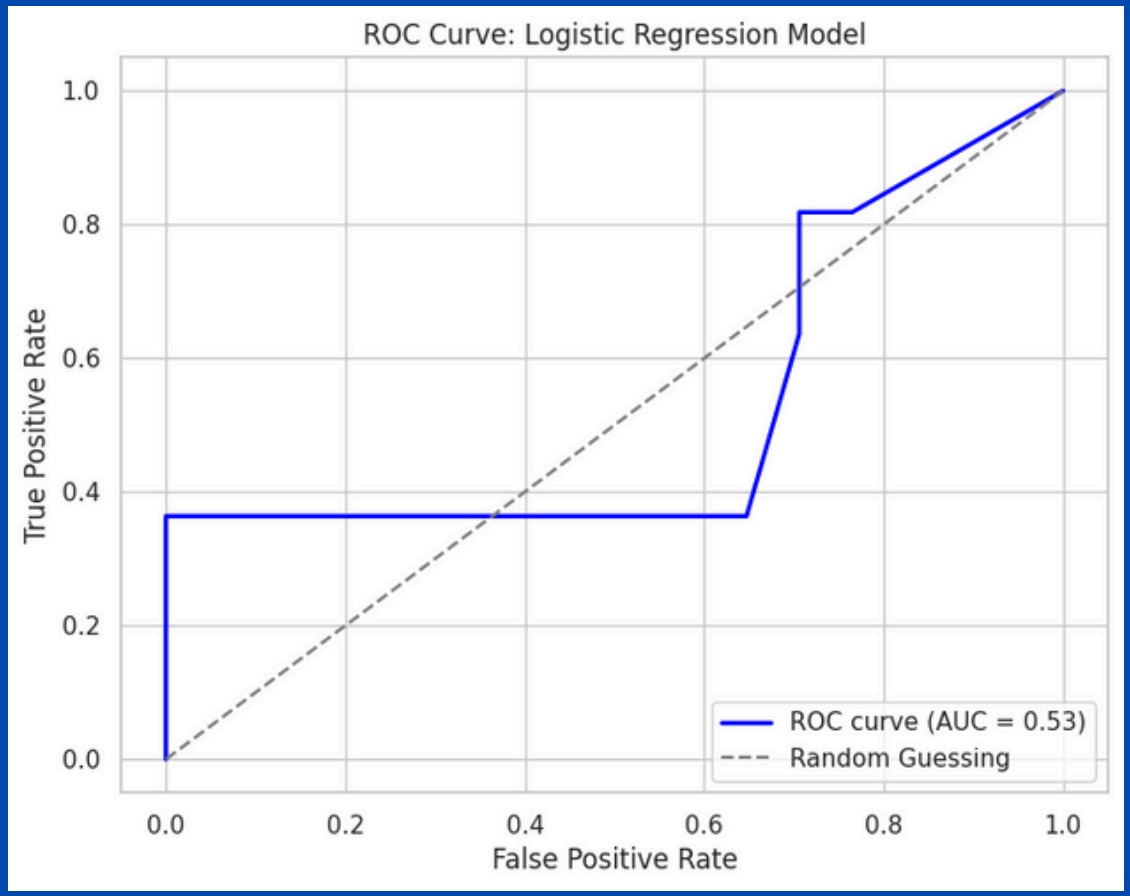

- MEASURES THE MODEL'S ABILITY TO SEPARATE WINS FROM LOSSES ACROSS ALL PROBABILITY THRESHOLDS.
- THE AUC $\approx$ 0.53, SLIGHTLY ABOVE RANDOM GUESSING, SHOWING THAT DISTINGUISHING WINS IS DIFFICULT WITH LIMITED FEATURES.
- REFLECTS THE UNPREDICTABLE NATURE OF NFL GAMES AND A MODESTLY PREDICTIVE MODEL.

Calibration Curve

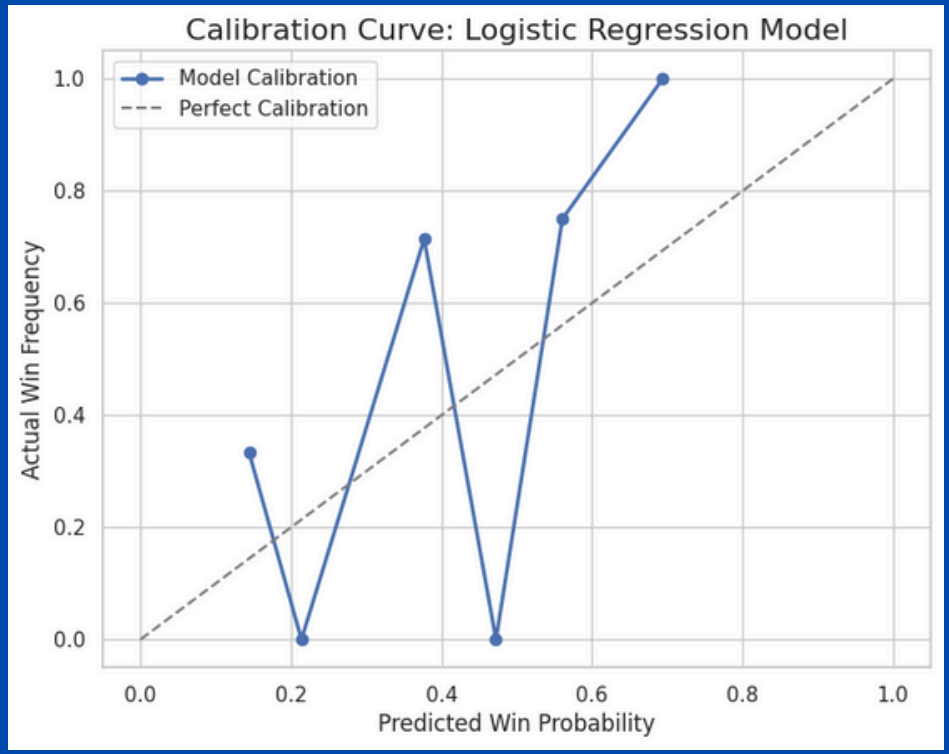

- COMPARES PREDICTED WIN PROBABILITIES WITH ACTUAL OUTCOMES IN THE TEST SET.
- IDEALLY, POINTS ALIGN WITH THE DIAGONAL LINE (PERFECT CALIBRATION).
- THE MODEL IS REASONABLY CALIBRATED AT LOWER PROBABILITY RANGES BUT BECOMES LESS RELIABLE AT HIGHER VALUES.
- THIS MATCHES EXPECTATIONS FOR LOGISTIC REGRESSION WITH LIMITED FEATURES.
- CONFIRMS THAT THE MODEL PRODUCES CAUTIOUS, REALISTIC PROBABILITY ESTIMATES RATHER THAN EXTREME PREDICTIONS.

2025 Season Prediction

OUTPUTS GENERATED:

- WIN PROBABILITY FOR EACH 2025 GAME
- PREDICTED WIN OR LOSS PER MATCHUP
- FINAL PROJECTED SEASON RECORD

FINAL PREDICTION: CHICAGO BEARS FINISH (5–12) IN 2025

WHY THIS RESULT OCCURS:

- MANY OPPONENTS HISTORICALLY OUTPERFORM THE BEARS.
- HOME FIELD ADVANTAGE PROVIDES ONLY MODEST IMPROVEMENT.
- MOST PREDICTED GAMES FALL INTO "CLOSE GAME" PROBABILITY RANGES (35%–55%).

Predicted Win Probability

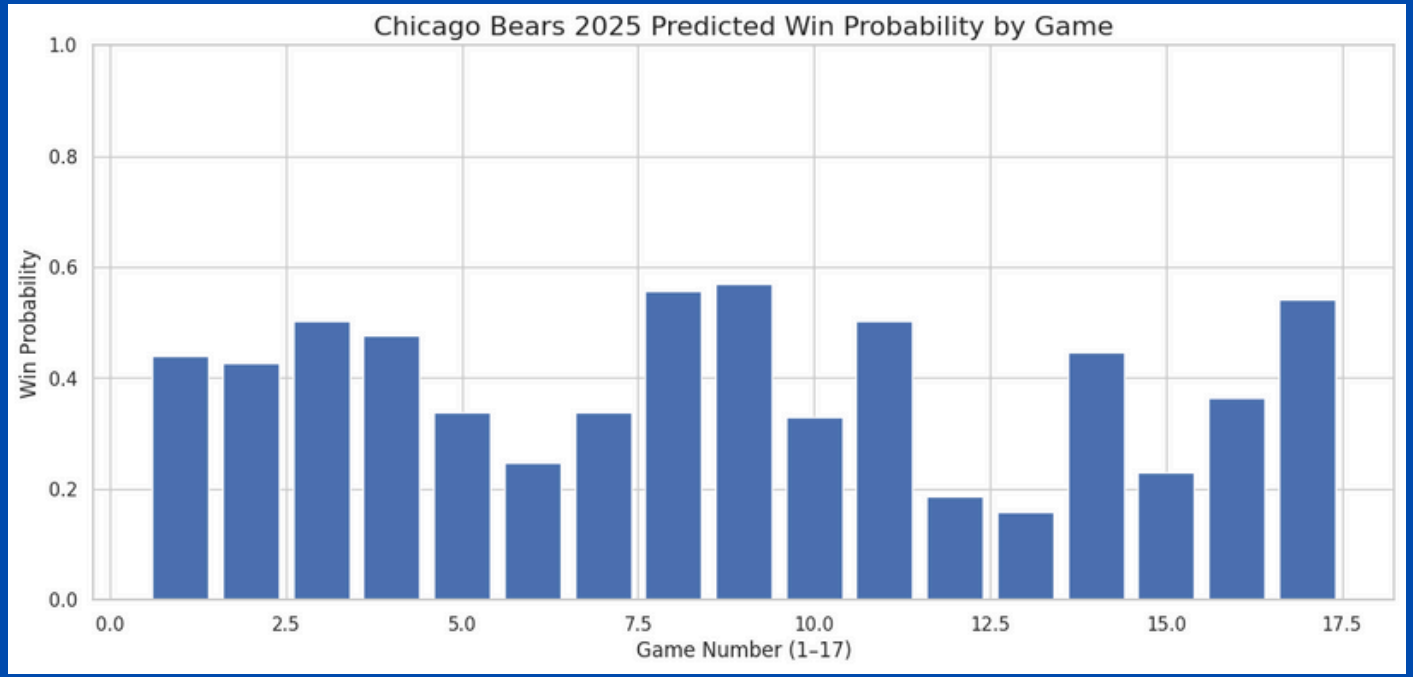

- Displays the model's predicted win probability for each of the 17 regular season games in 2025.
- Most predictions fall between 35% and 55%, indicating many games are statistically "close."
- Only a small number of games show probabilities above 55%, reflecting limited high confidence matchups.
- This pattern suggests a challenging schedule and emphasizes the Bears' inconsistency in recent years.

2025 Pie Chart

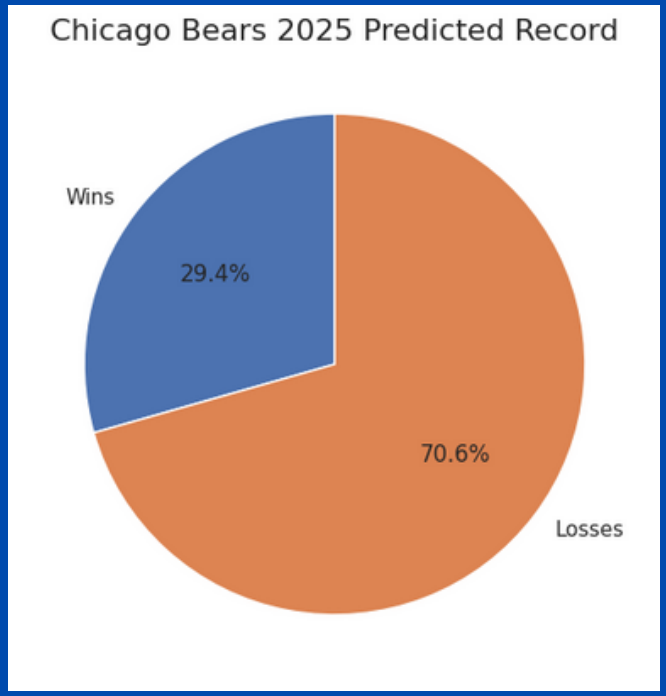

- THE OVERALL PROJECTED SEASON OUTCOME: 5 WINS AND 12 LOSSES.
- THE WIN PORTION HIGHLIGHTS THE BEARS' EXPECTED STRUGGLES IN 2025.
- THE PIE CHART COMMUNICATES THE BROADER SEASON OUTLOOK IN A SIMPLE, INTUITIVE WAY.

Win Probability Distribution

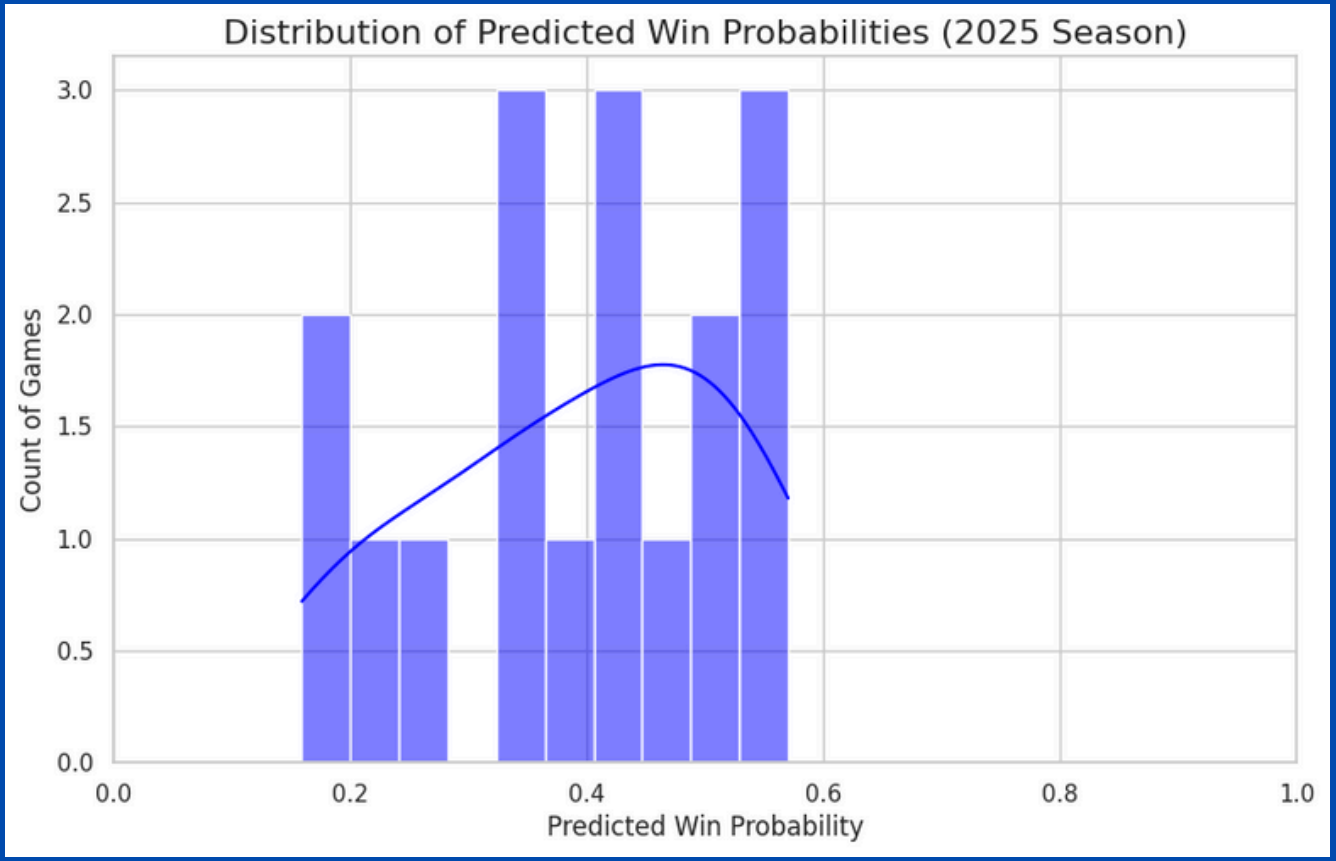

- Displays the distribution of win probabilities across all 2025 games.
- Shows that predictions generally cluster around mid-range probabilities, rather than extreme values.
- Indicates the model avoids overconfident predictions appropriate given the simple feature set.
- The distribution reveals that the Bears are rarely strong favorites, consistent with the final 5–12 projection.

Home vs Away Win Probability

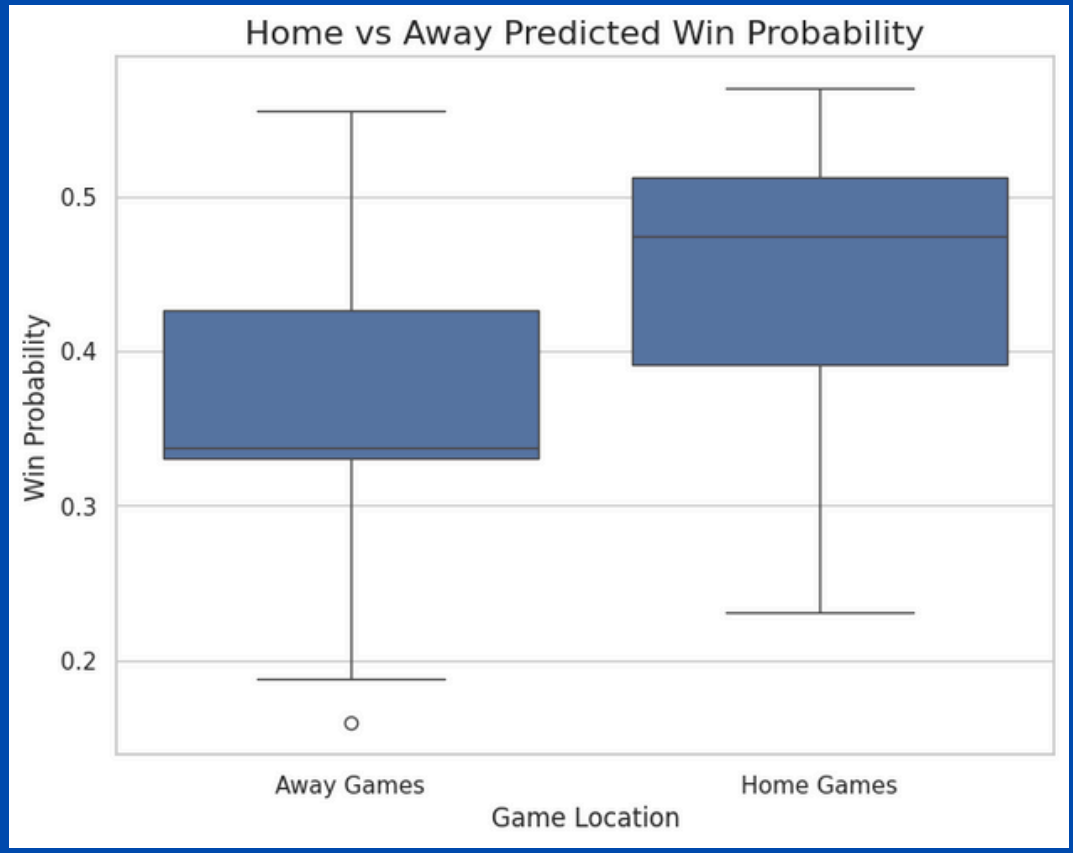

- Shows that the Bears' predicted win probability is higher at home than away, consistent with historical home-field advantage.
- Home games have a higher median probability and a more compact spread.
- Away games show lower median performance with greater variability.
- Confirms the model successfully captured one of the strongest contextual predictors in the NFL: game location.

Conclusion

What this project demonstrates

- Data science execution: cleaning → modeling → visualization
- Ability to work with real sports datasets
- Competency in supervised learning (logistic regression)
- Model interpretability and evaluation of uncertainty

What the model predicts

- A realistic 5–12 record for the Bears
- Home-field advantage increases win probability
- Opponent strength is the strongest predictor in the dataset

Limitations

- Only uses two core features (simple, interpretable model)
- NFL outcomes are highly variable, with random events not captured by past data
- More granular data (player stats, injuries, play-by-play) could improve accuracy

Future work (if expanded into full capstone)

- Incorporate more advanced features (QB stats, team Elo, injuries, weather)
- Compare logistic regression to tree-based models or neural nets
- Build an interactive dashboard for exploring predictions